

Google™



Building high-throughput data pipelines on Google App Engine

Brett Slatkin
May 20th, 2010

**View live notes and ask questions about
this session on Google Wave**

<http://tinyurl.com/app-engine-pipelines>

Me

<http://onebigfluke.com>

Agenda

- Intro
- Fan-out
- Transactional sequences
- Fan-in
- Bonus round
- Future directions

Intro



What are pipelines?

- Constant trickle/torrent of inputs and outputs
 - Assembly-line
- Optimize for end-to-end latency of input to output (~seconds)
- Minimize incremental cost
- Not lossy, eventually consistent, all inputs served

What are **NOT** pipelines?

- Offline systems like MapReduce
- Batch processing, report generation
- Outputs are from a snapshot of inputs
- Latency from input to output is ~hours

Example apps

- Pipelines

- Email, Twitter, PubSubHubbub (routing)
- Reddit, Digg (voting, agg)
- CRM (~yeah, really)

- Not pipelines

- Guestbook (flat)
- Terasort (snapshot)
- Chat (transient)

- Hybrid

- YouTube, Vimeo (transcode)
- Flickr, Picasa (face recog, tags)

Fan-out: Continuations

What is fan-out?

- One action leads to many others
- Datastore-based inbox systems (eg, microblogging)
- Send notification emails, XMPP, SMS, Channel API, APN
- Web service calls
- Enqueue more tasks

Example fan-out

- Update a party invitation, send an email to everyone


```
class Party(db.Model):  
    when = db.DateTimeProperty()  
    host = db.UserProperty()  
class PartyGoer(db.Model):  
    party = db.ReferenceProperty(Party)  
    name = db.StringProperty()  
    address = db.EmailProperty()
```


Continuation passing (naively)

```
class EmailHandler(webapp.RequestHandler):
    def post(self):
        my_party = self.request.get("party_key")
        cursor = self.request.get("cursor")
        query = PartyGoer.all().filter(
            "party =", db.Key(my_party))
        if cursor:
            query.with_cursor(cursor)
        goers = query.fetch(10)
        # Send some emails ...
        if len(goers) == 10:
            taskqueue.add(url='/work/email',
                params={'party_key': my_party,
                    'cursor': query.cursor()})
```

Continuation passing (the wrong way)


- Any failures and...



Continuation passing (the right way)

```
class EmailHandler(webapp.RequestHandler):
    def post(self):
        my_party = self.request.get("party_key")
        cursor = self.request.get("cursor")
        query = PartyGoer.all().filter(...)
        if cursor:
            query.with_cursor(cursor)
        goers = query.fetch(10)
        if len(goers) == 10:
            taskqueue.add(
                url='/work/email',
                params={'party_key': my_party,
                       'cursor': query.cursor()},
                name=int(self.request.get('gen')) + 1)
        # Send some emails ...
```

Continuation passing (the right way)



Continuation passing benefits

- Failures and spurious retries are isolated
- Execute continued work in parallel

Continuation passing benefits 2

- Pairs well with asynchronous APIs
 - Async URLFetch in Python
 - Java support since 1.3.1 (February)
 - Async Datastore
 - Python: <http://asynctools.googlecode.com>
 - Java: <http://twig-persist.googlecode.com>
- Used in PubSubHubbub reference hub
 - **100-300** worker requests/sec constantly

Transactional sequences

What are transactional sequences?

- Datastore transactions and transactional tasks
- Guarantee that tasks run after data is written
 - Strong consistency when task is run
- Enables roll-forward semantics to fanned-out data
 - Build materialized views

What are materialized views good for?

- A query that's saved back into the database
 - Read-heavy, cached, secondary indexes
 - Eventually consistent views
- Incremental aggregations (commutative)
- Natural and left-joins
- Filter/query/sorting materialized results

SQL Example: Students in school

Student	Grade
Bob	4
Daisy	3
...	

```
SELECT grade, count(*) as count
FROM Student
GROUP BY grade;
```

grade	count
3	5
4	7
...	

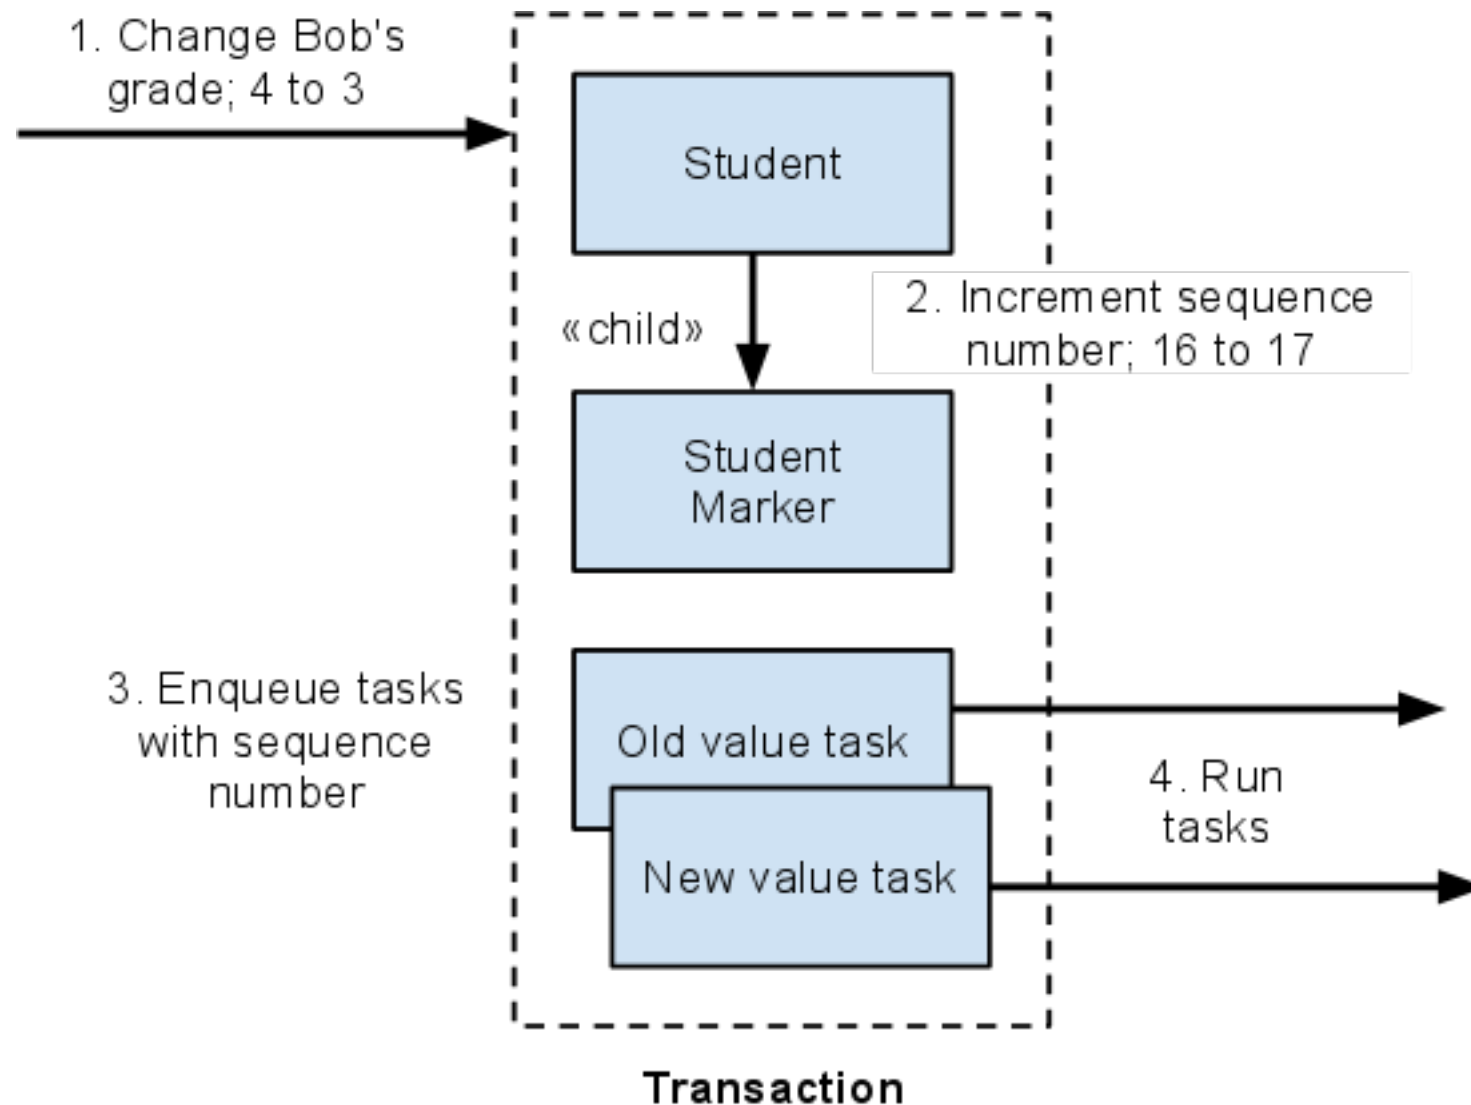
App Engine Example: Students in school

```
class Student(db.Model):  
    name = db.StringProperty()  
    grade = db.IntegerProperty()
```

```
class Marker(db.Model):  
    sequence = db.IntegerProperty(default=0)  
    present = db.BooleanProperty()
```

```
class GroupCount(db.Model):  
    grade = db.IntegerProperty()  
    count = db.IntegerProperty(default=0)
```

Roll-forward semantics: Update source data



Update source data

```
def update(name, new, id):  
    def txn():  
        if id:  
            student = Student.get_by_id(id)  
            old, student.grade = student.grade, new  
        else:  
            student = Student(name=name, grade=new)  
            student.put() # Assign ID  
            old, id = None, student.key().id()  
  
    marker_key = db.Key.from_path(  
        'Marker', id, parent=student.key())  
    marker = db.get(marker_key)  
    if not marker: marker = Marker(key=marker_key)  
    marker.sequence += 1  
    # continues on next slide
```

Update source data continued

```
db.put([student, marker])
taskqueue.Task(
    url='/work',
    params={'student_id': id, 'grade': new,
           'sequence': marker.sequence,
           'present': True}
).add(transactional=True)
if old is not None:
    taskqueue.Task(
        url='/work',
        params={'student_id': id, 'grade': old,
               'sequence': marker.sequence,
               'present': False}
    ).add(transactional=True)
db.run_in_transaction(txn)
```

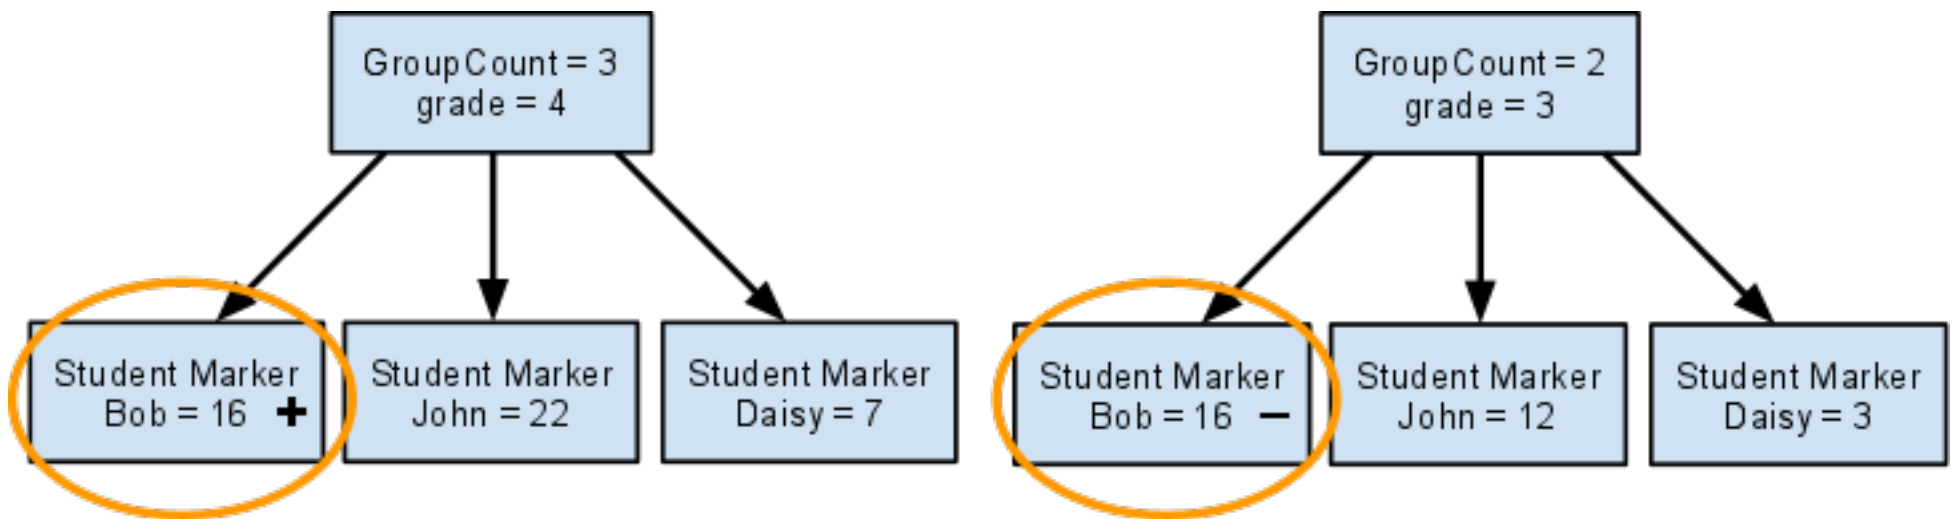
App Engine Example: Students in school

```
class Student(db.Model):  
    name = db.StringProperty()  
    grade = db.IntegerProperty()
```


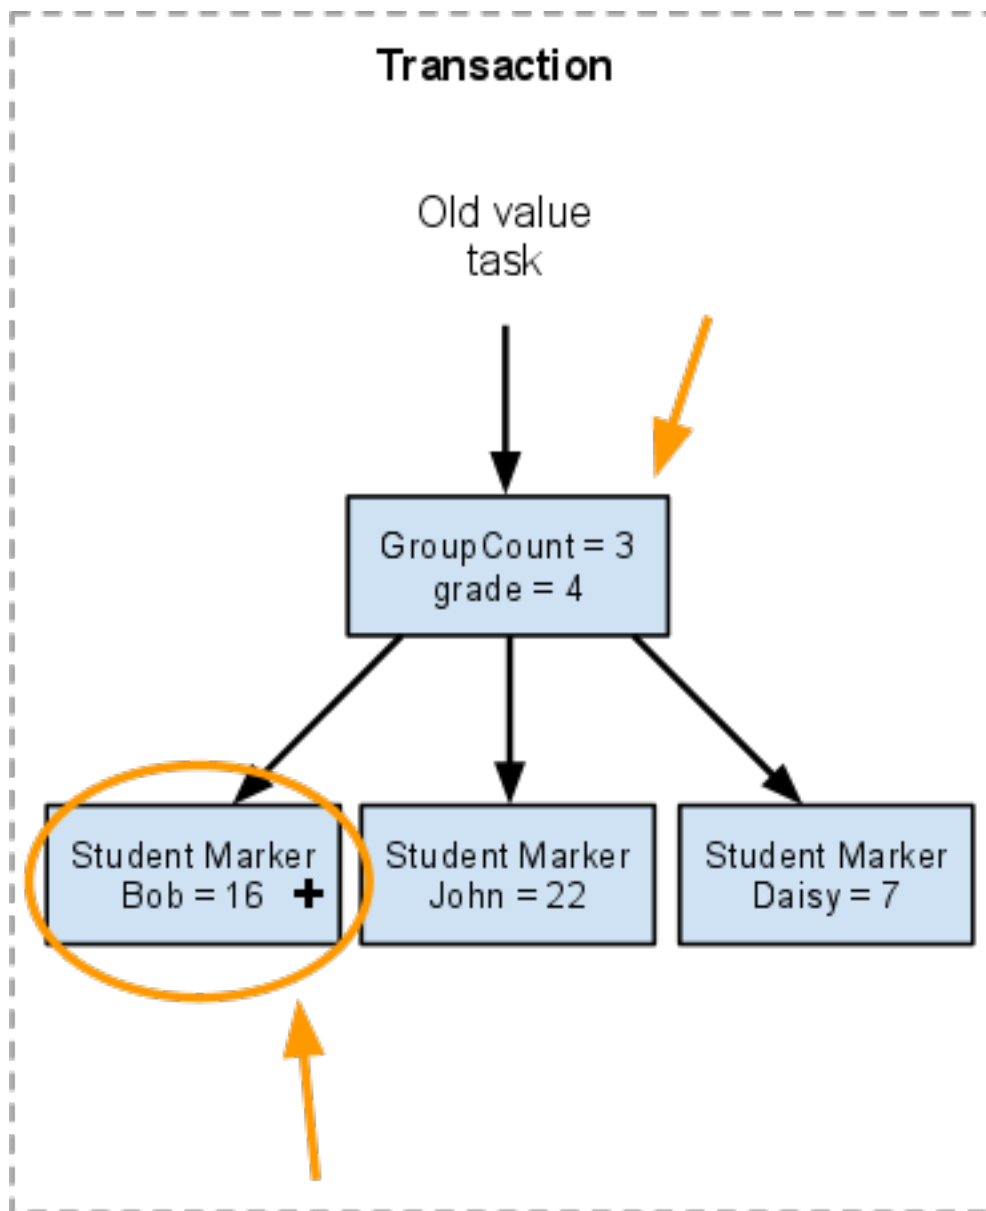
```
class Marker(db.Model):  
    sequence = db.IntegerProperty(default=0)  
    present = db.BooleanProperty()
```

```
class GroupCount(db.Model):  
    grade = db.IntegerProperty()  
    count = db.IntegerProperty(default=0)
```

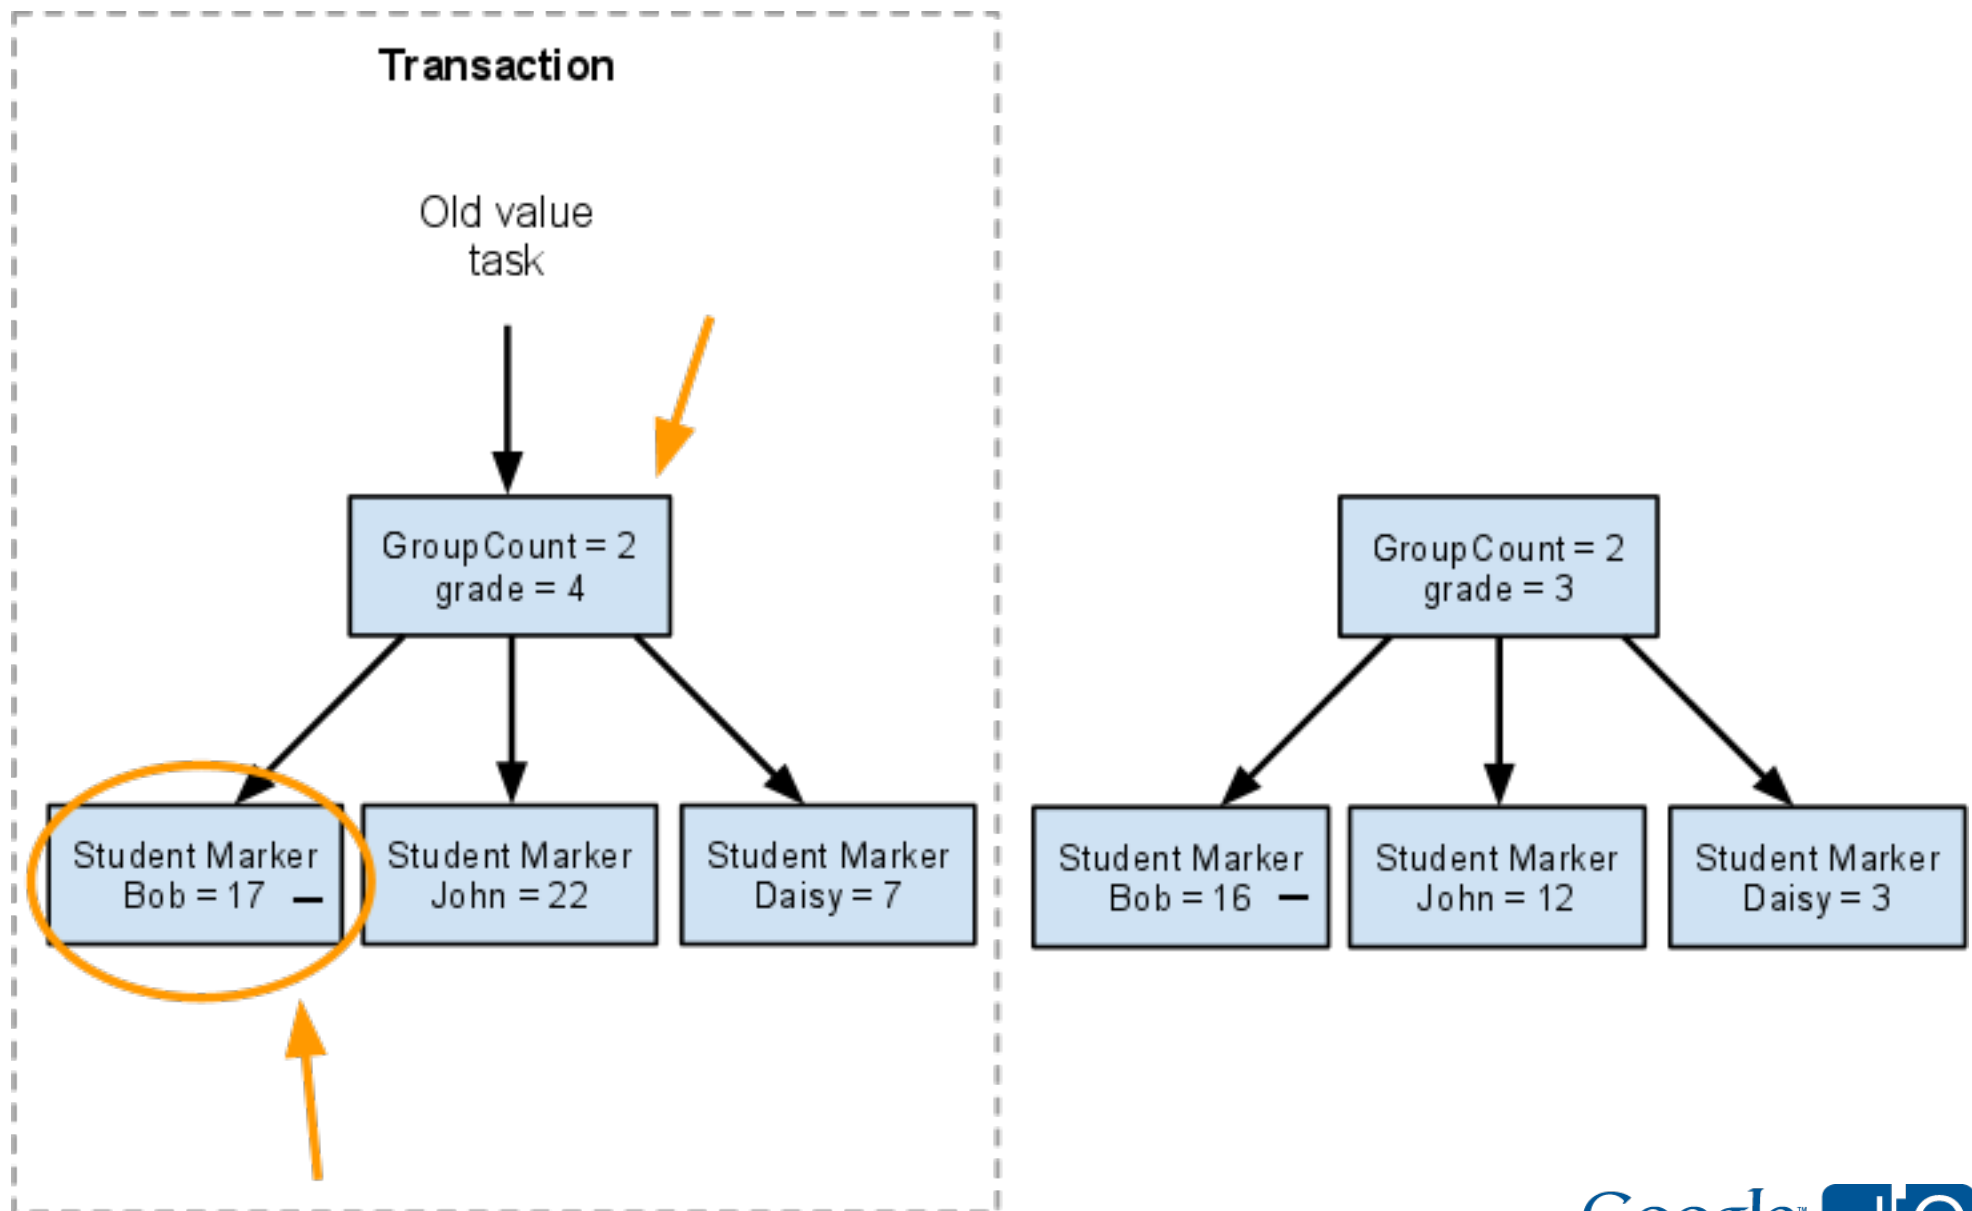
Roll-forward semantics: View initial state



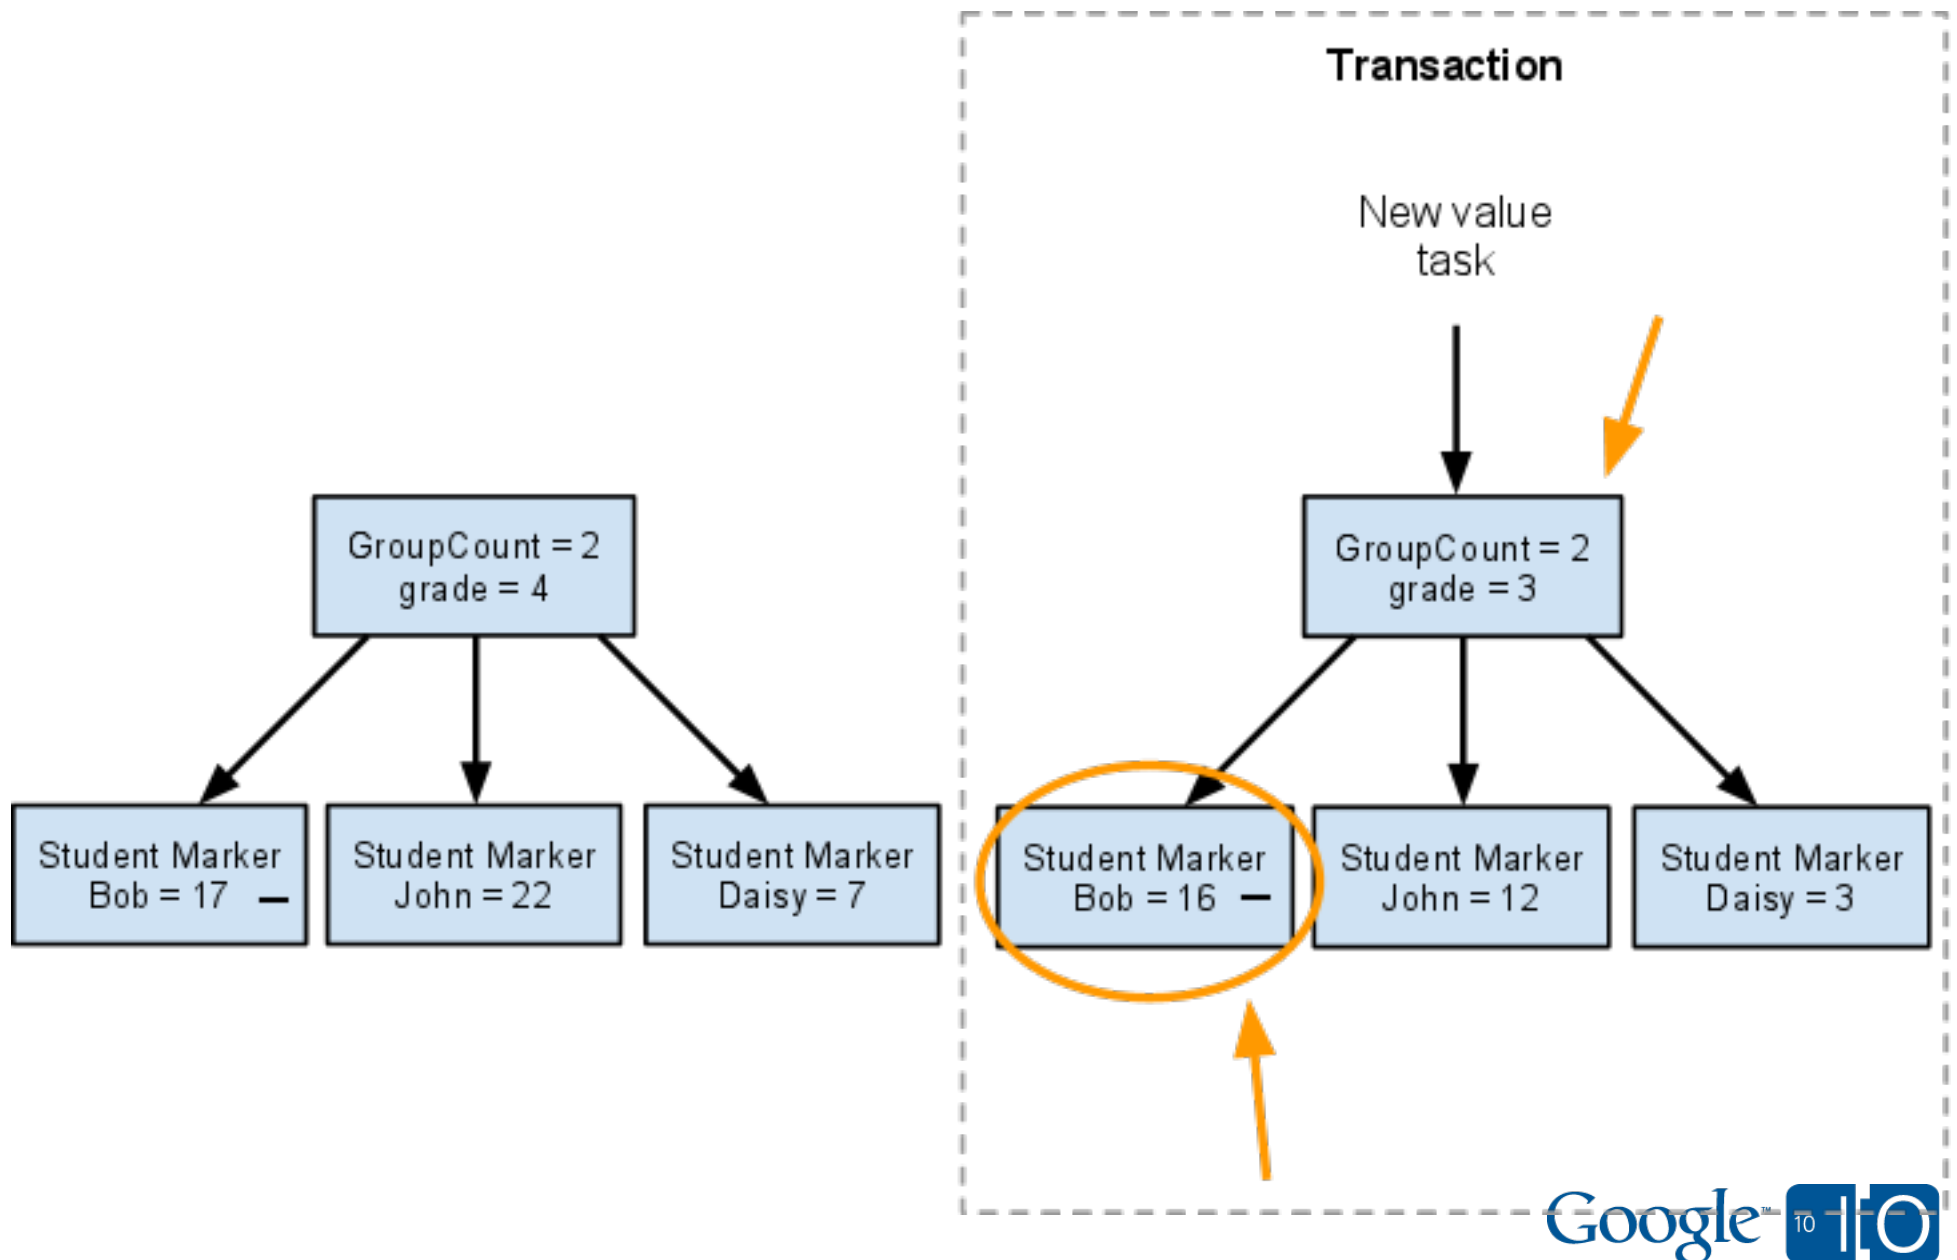
Roll-forward semantics: Update old value



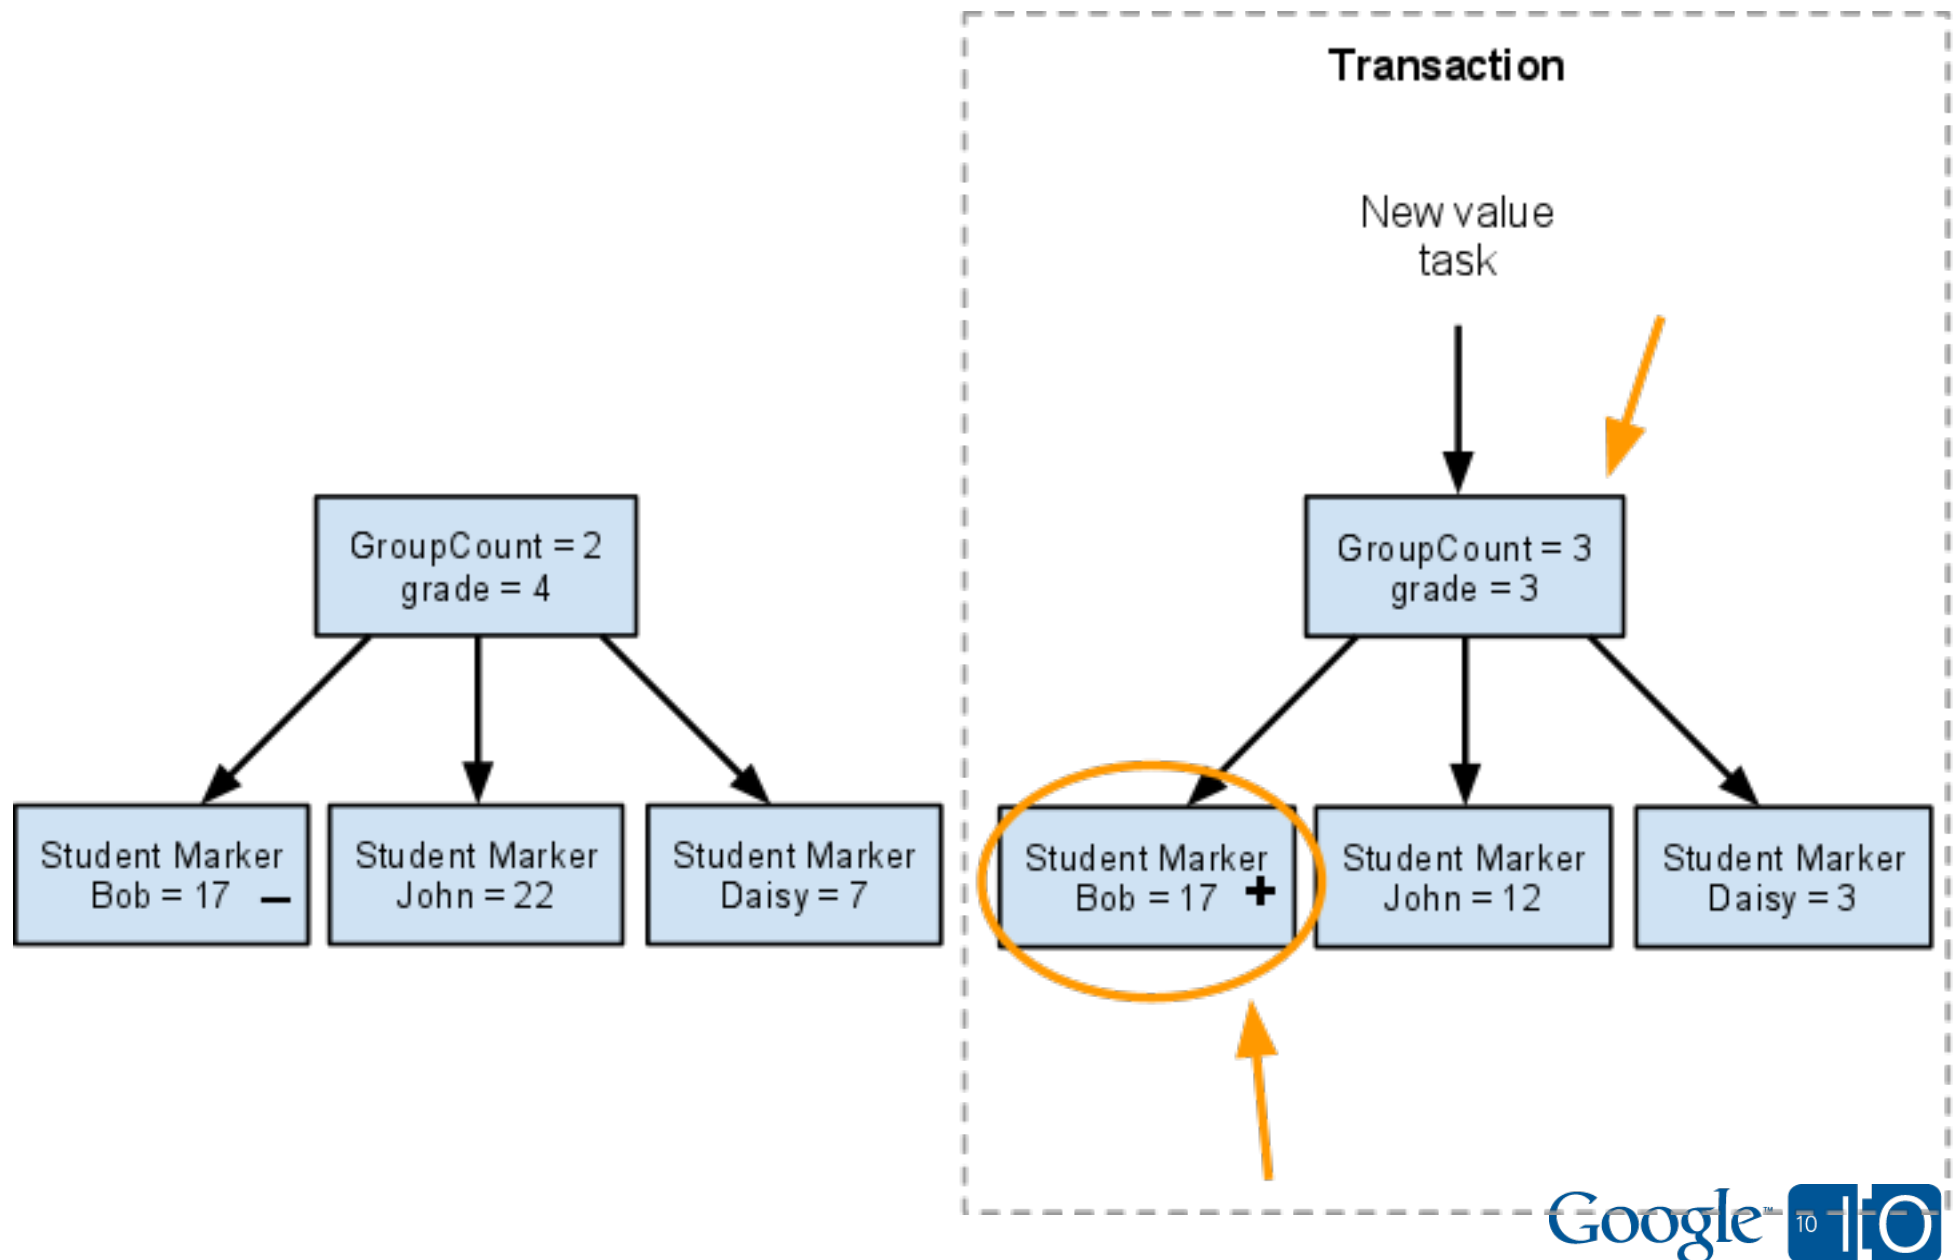
Roll-forward semantics: Update old value



Roll-forward semantics: Update new value



Roll-forward semantics: Update new value



Update group counts

```
def apply(sequence, present, grade, id)
    group_key = db.Key.from_path('GroupCount', grade)
    marker_key = db.Key.from_path(
        'Marker', id, parent=group_key)

def txn():
    group, marker = db.get([group_key, marker_key])
    if not group:
        group = GroupCount(key=group_key)
    if not marker:
        marker = Marker(key=marker_key)
    if marker.sequence >= sequence:
        raise db.Rollback('Ignore out-of-order')
    # continues on next slide
```

Update group counts continued

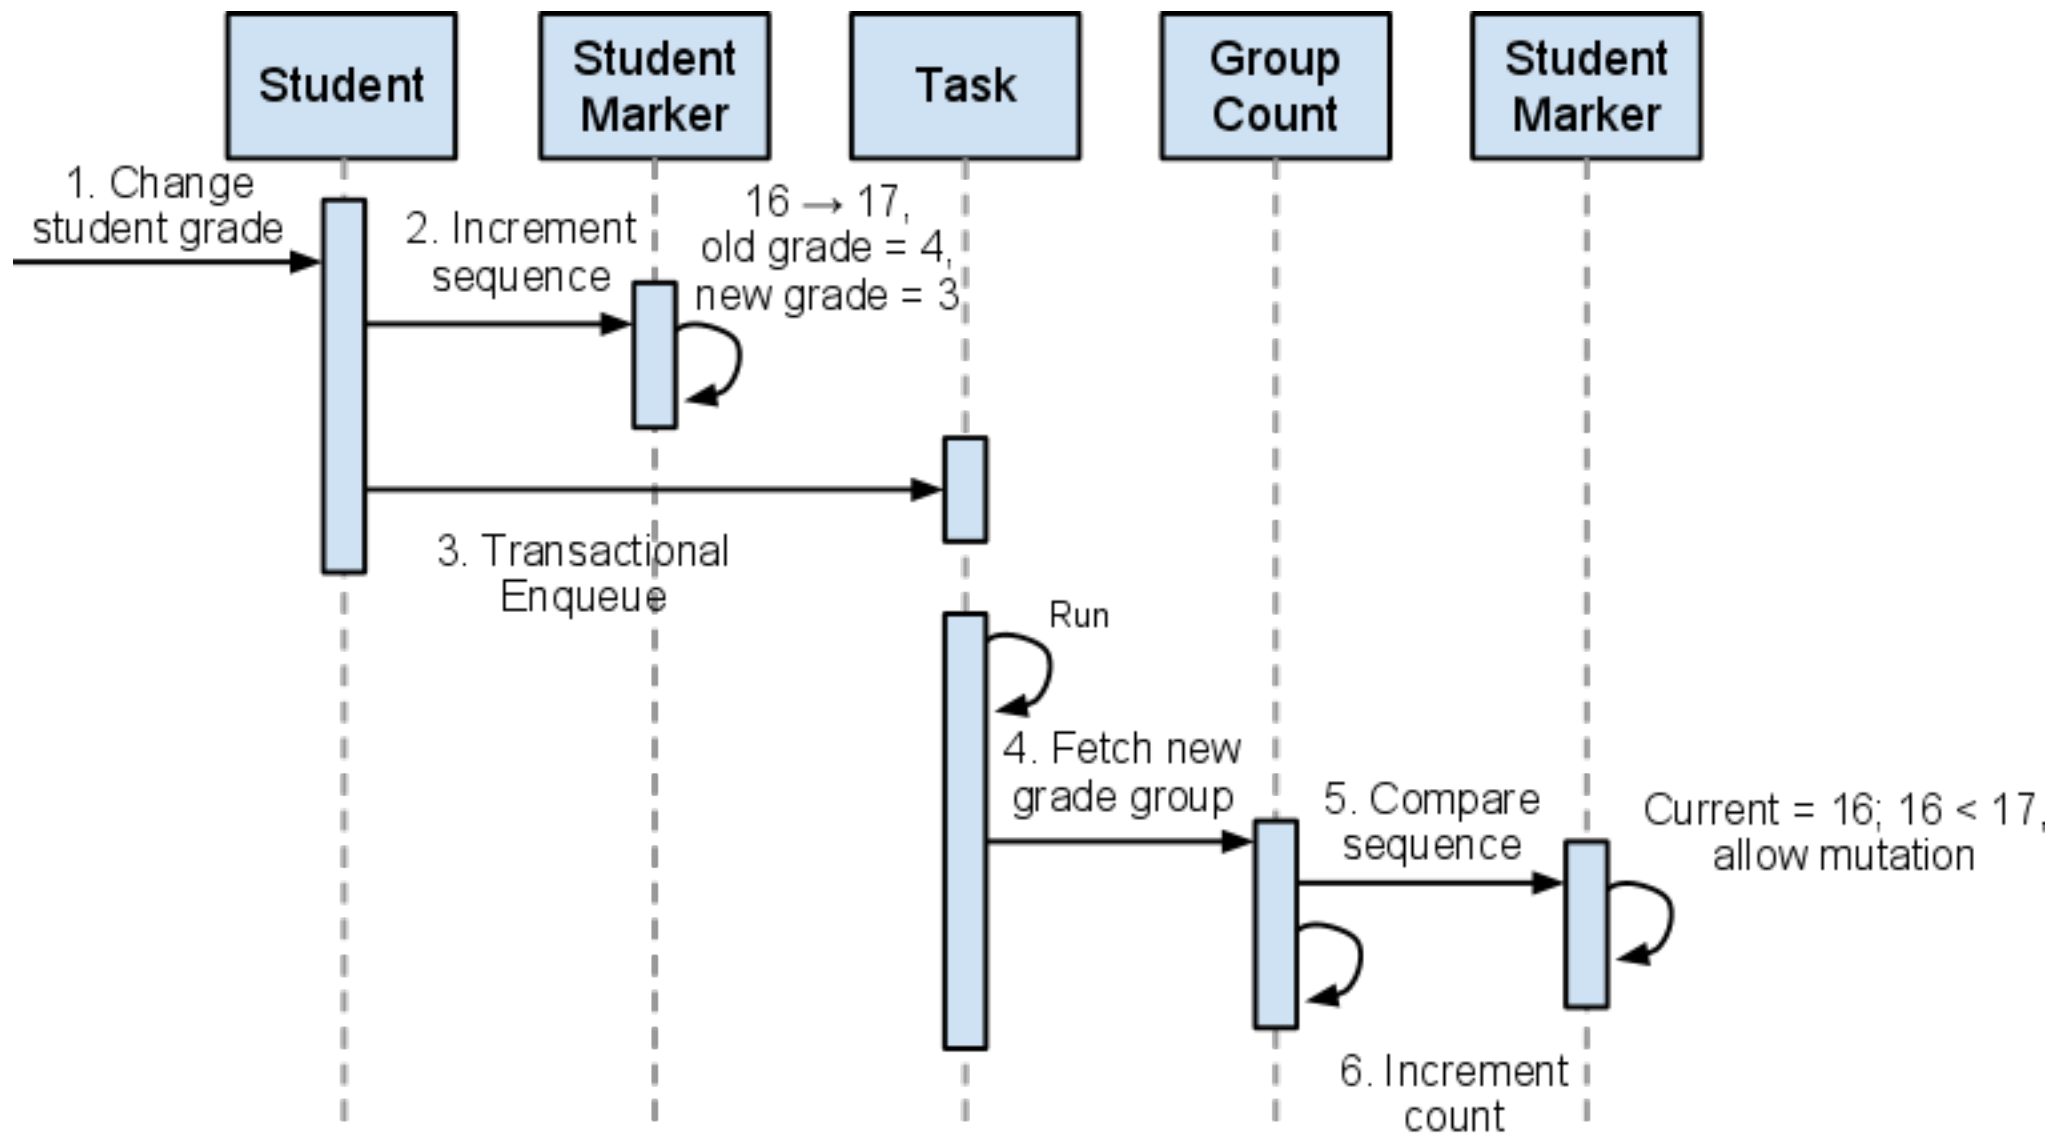
```
old, marker.present = marker.present, present
marker.sequence = sequence
db.put(marker)
```

```
if old:
    group.count -= 1
if present:
    group.count += 1
```

```
if group.count == 0:
    group.delete()
else:
    db.put(group)
```

```
db.run_in_transaction(txn)
```

Coordination sequence



Transactional sequences demo

Sequencing details

- Each aggregation row is in its own entity group
 - Update aggregation rows in separate transactions
 - Order of task application doesn't matter
- Marker entity is child of each Count (aggregation) row
 - Marker indicates presence of Student in aggregation
 - Sequence numbers let you ignore old/stale updates
- **Bridge transactions** across entity groups

Sequencing details 2

- Works well for commutative operations (count, sum)
 - Toggling presence is add or subtract
 - Ancestor queries for more complex functions
 - Use continuations and cursors to continue queries
- Enqueue multiple tasks at source data write time
 - Update many aggregations in parallel

Sequencing details 3

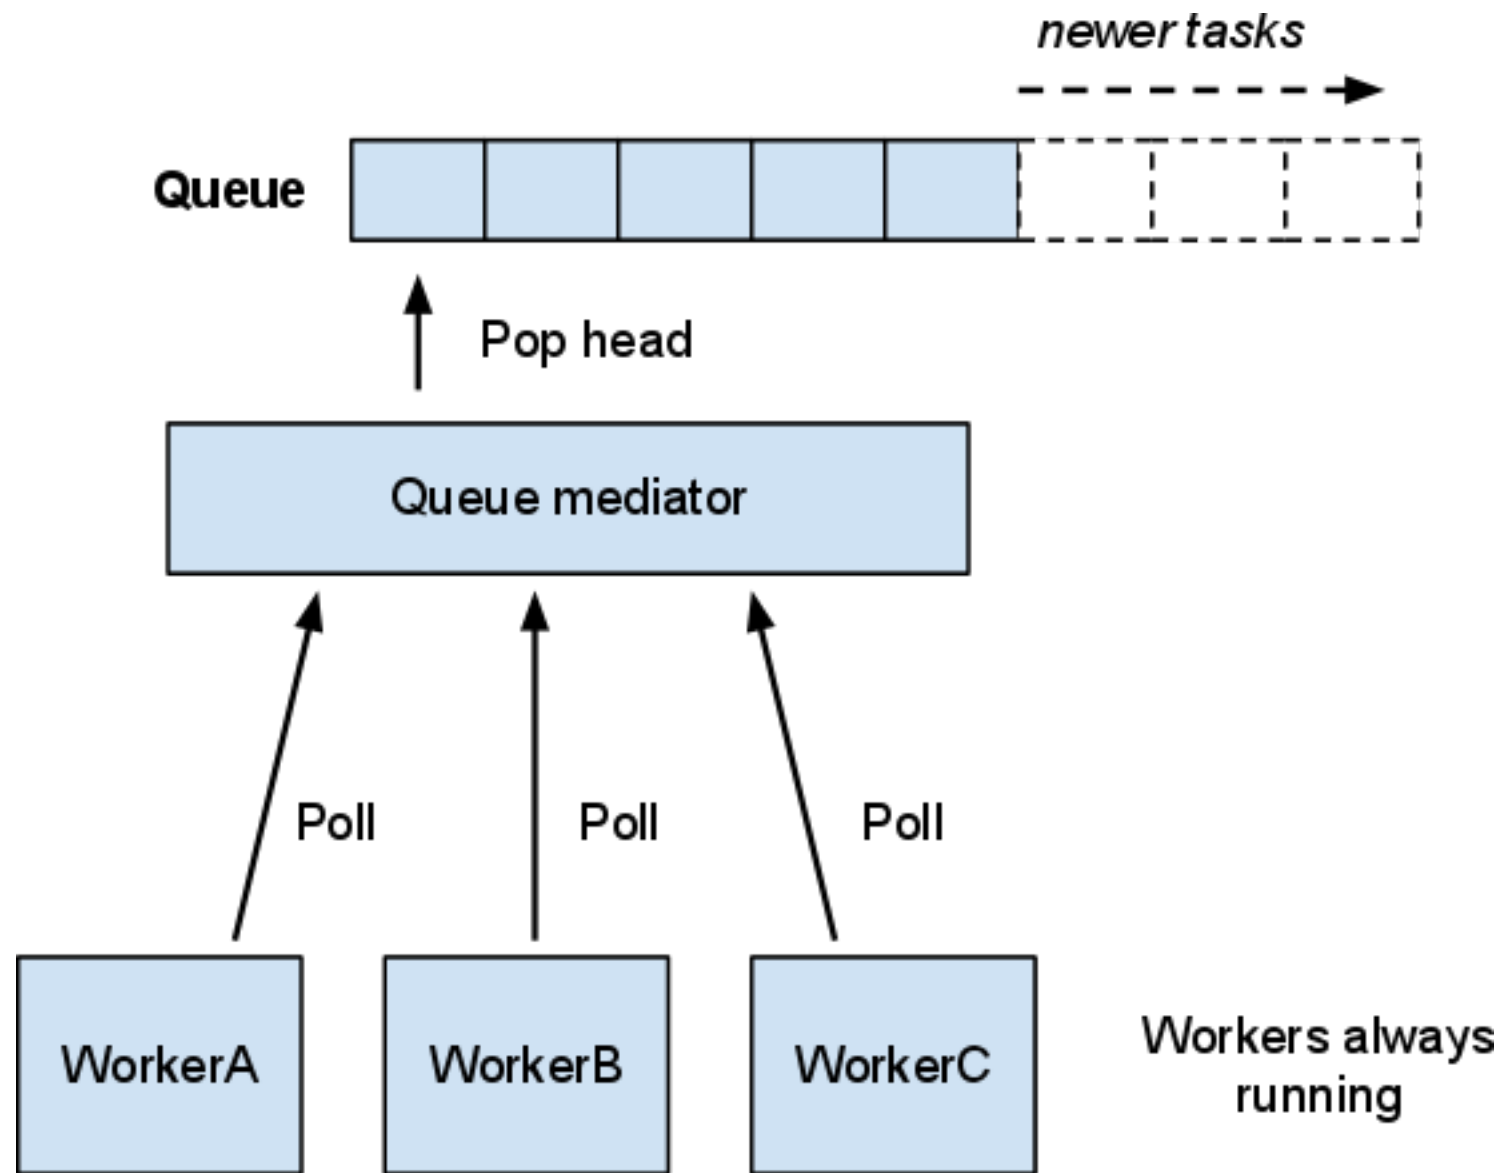
- Max throughput proportional to number of aggregation rows
 - Watch out: Data distribution across aggregation
 - Rate-limit materialized view tasks to 1/sec for safety
- Storage cost for "presence" entities

Fan-in: Fork-join queues

Why fan in?

- Apply multiple data transforms in batches
 - Counters, aggregations, roll-ups, reservations
 - Reddit/Digg-style: Save users' voting history
 - Beat the ~ 1 write/sec per entity group safety margin
- Wait for high-latency API calls simultaneously
 - RSS aggregators, microblog data sinks
 - Use fewer threads = more throughput
 - Ensure queues do not back up
- Amortize overhead costs with parallel work

Polling workers: Traditional approach



Polling workers: Problems

- How many polling workers do you need to ensure 50 new tasks per second are serviced within 500ms?
- What if tasks take at least 10 seconds?
- How do you guarantee exclusivity?


Polling workers are for offline processing

- Not dynamic, not low-latency
- Must have enough workers running to match peak load
- Eliminates the benefits of App Engine's *push* task queue

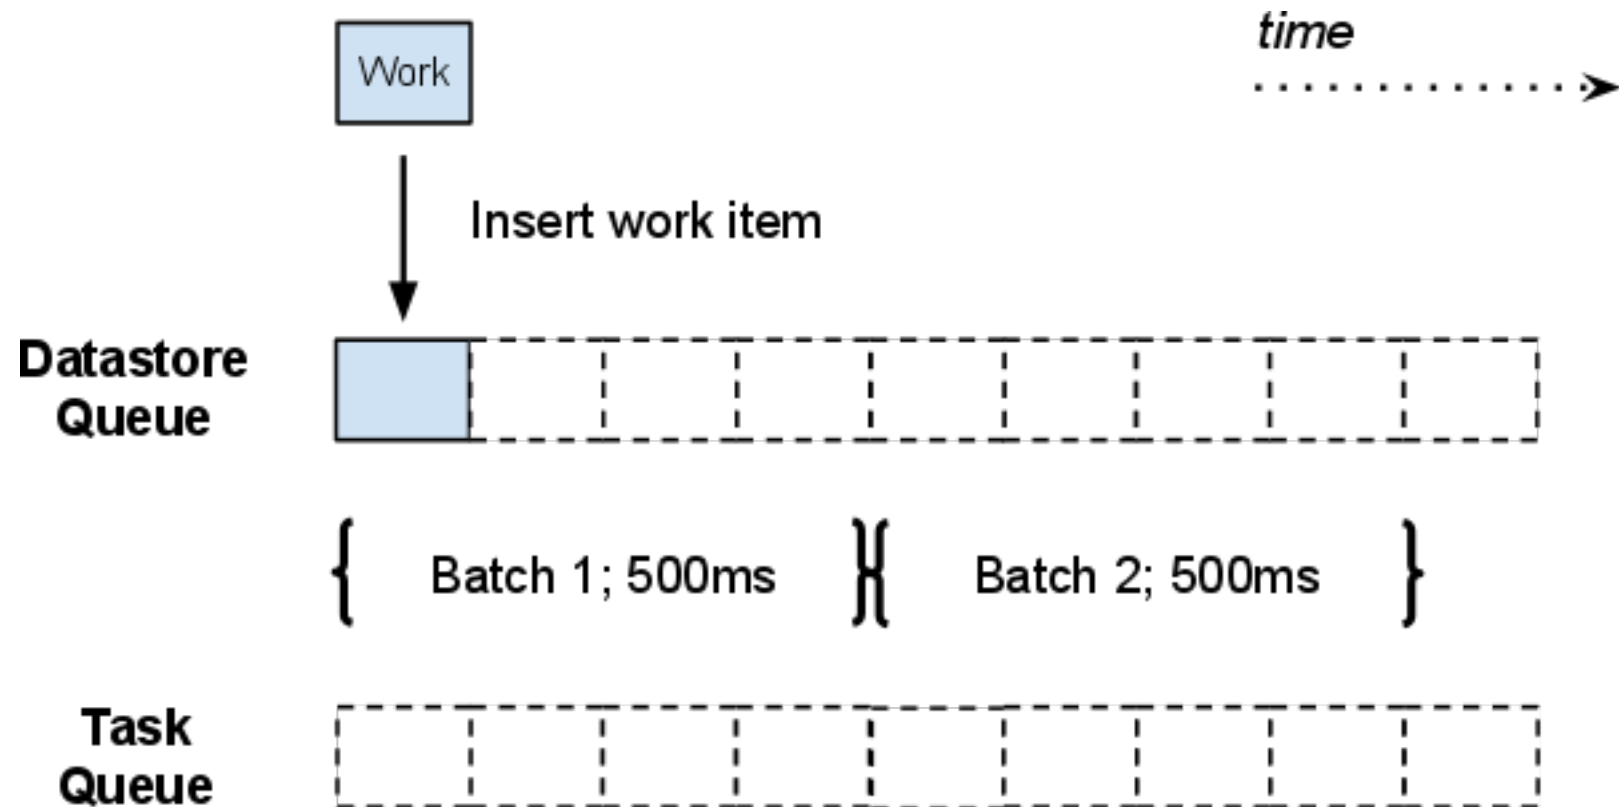
What is a fork-join queue?

- Fork incoming work items apart as they enter
- Work starts within maximum fixed period after arrival
- Execute work in batch for efficiency
- Join completed work items together into a result (optional)


Fork-join queue with Datastore



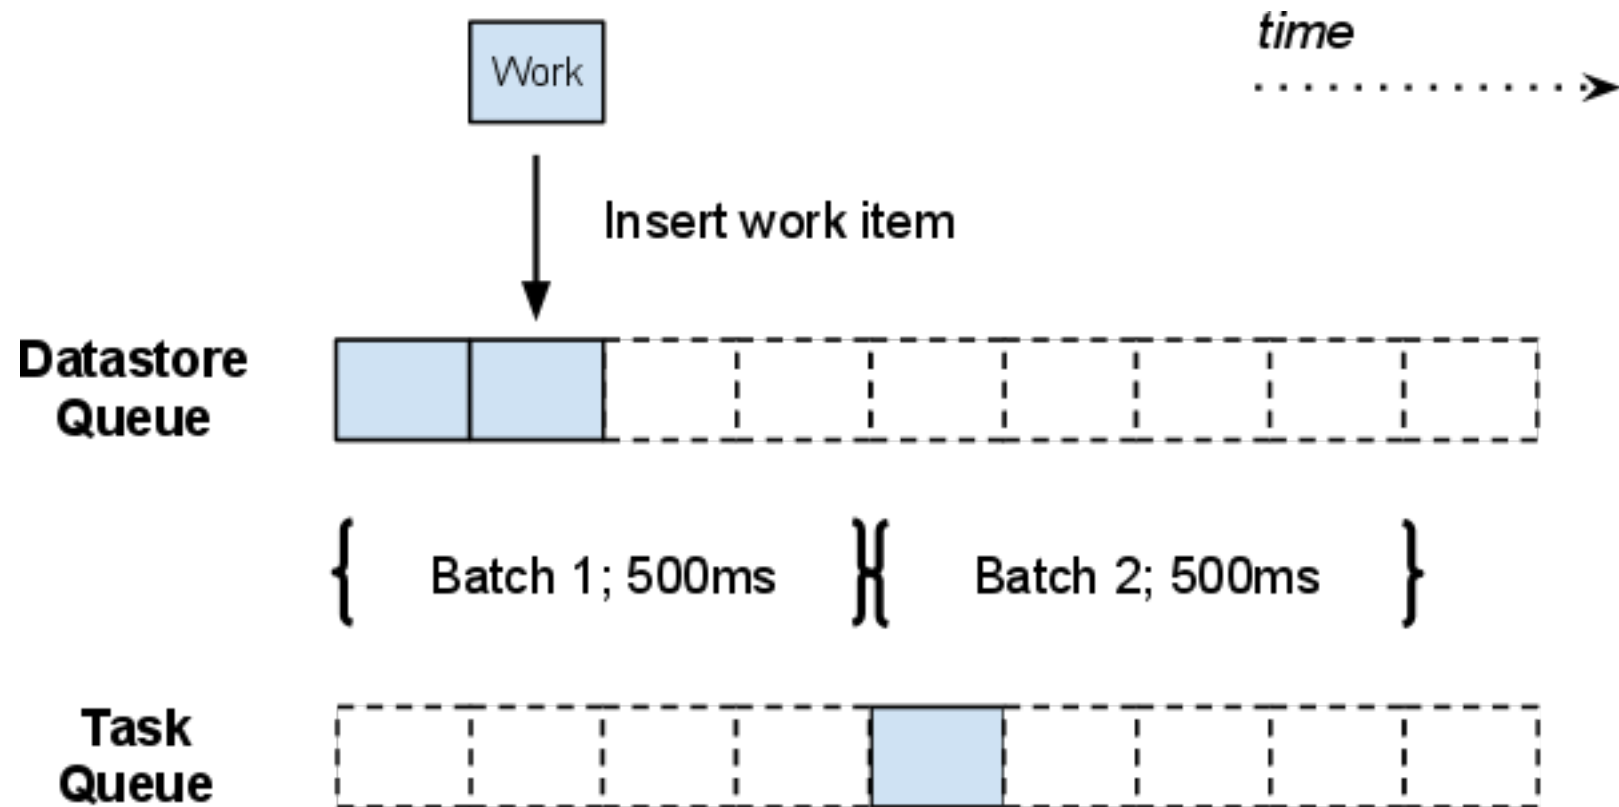
Fork-join queue with Datastore



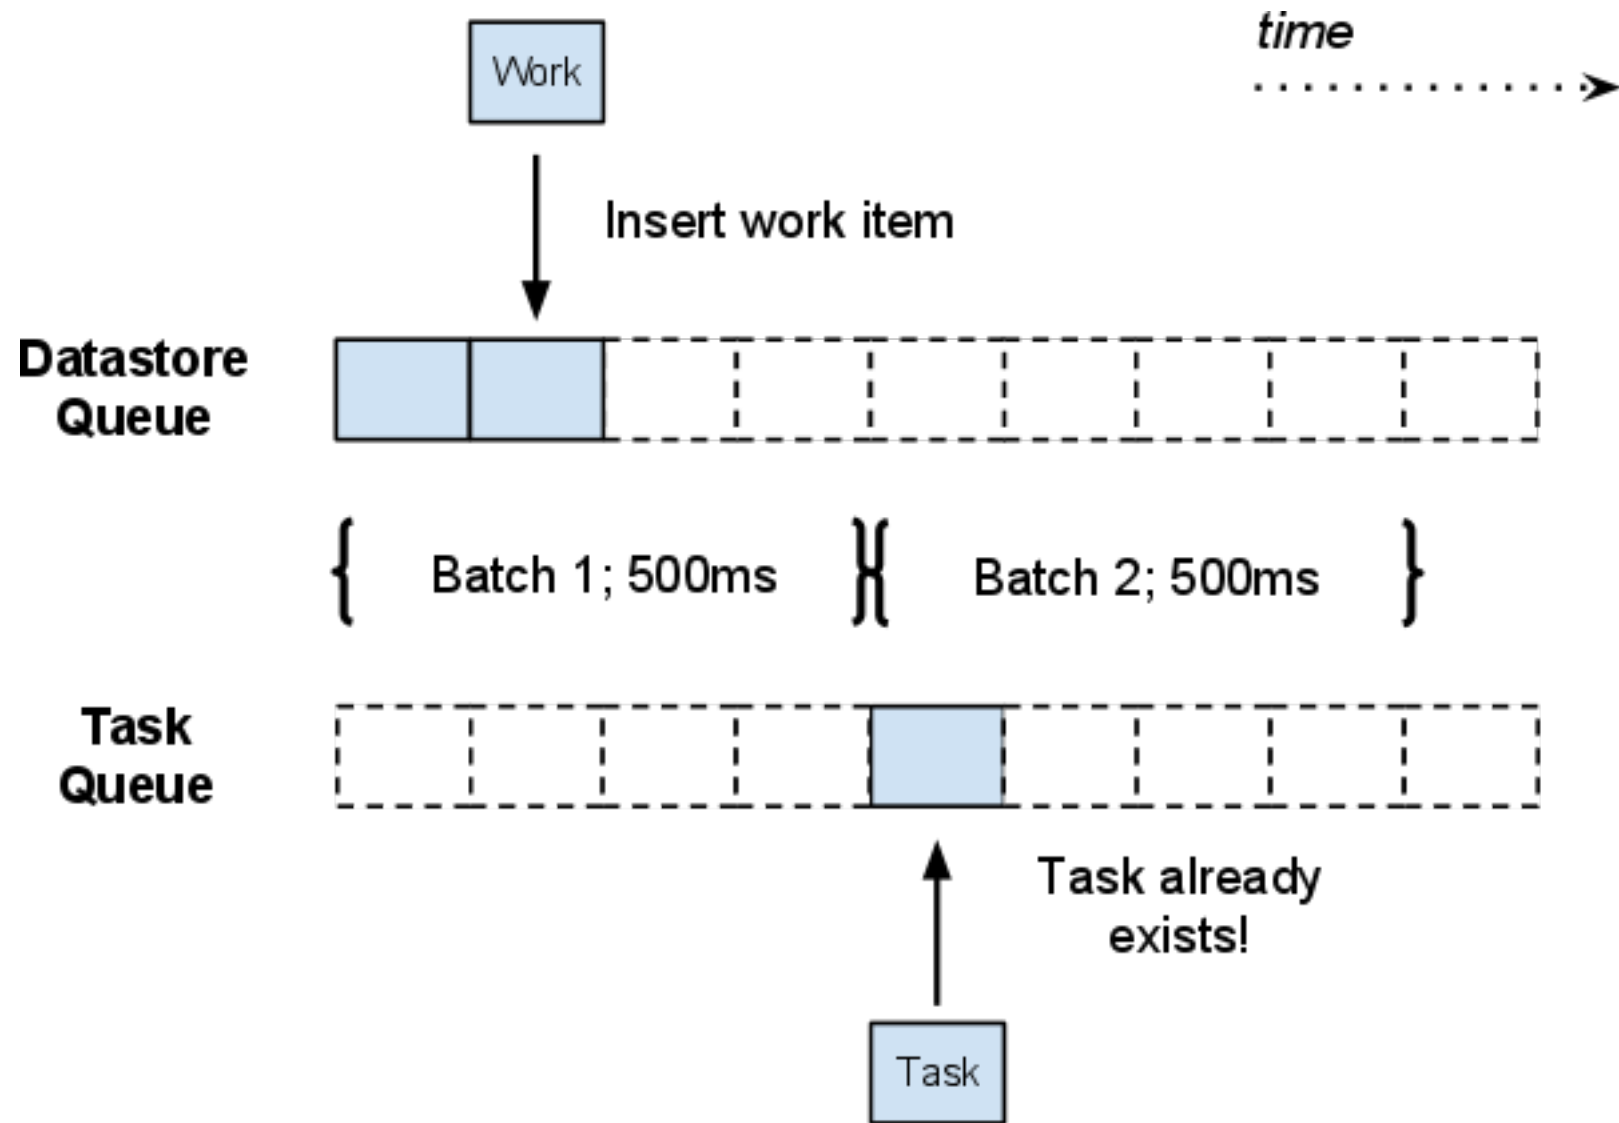
Fork-join queue with Datastore



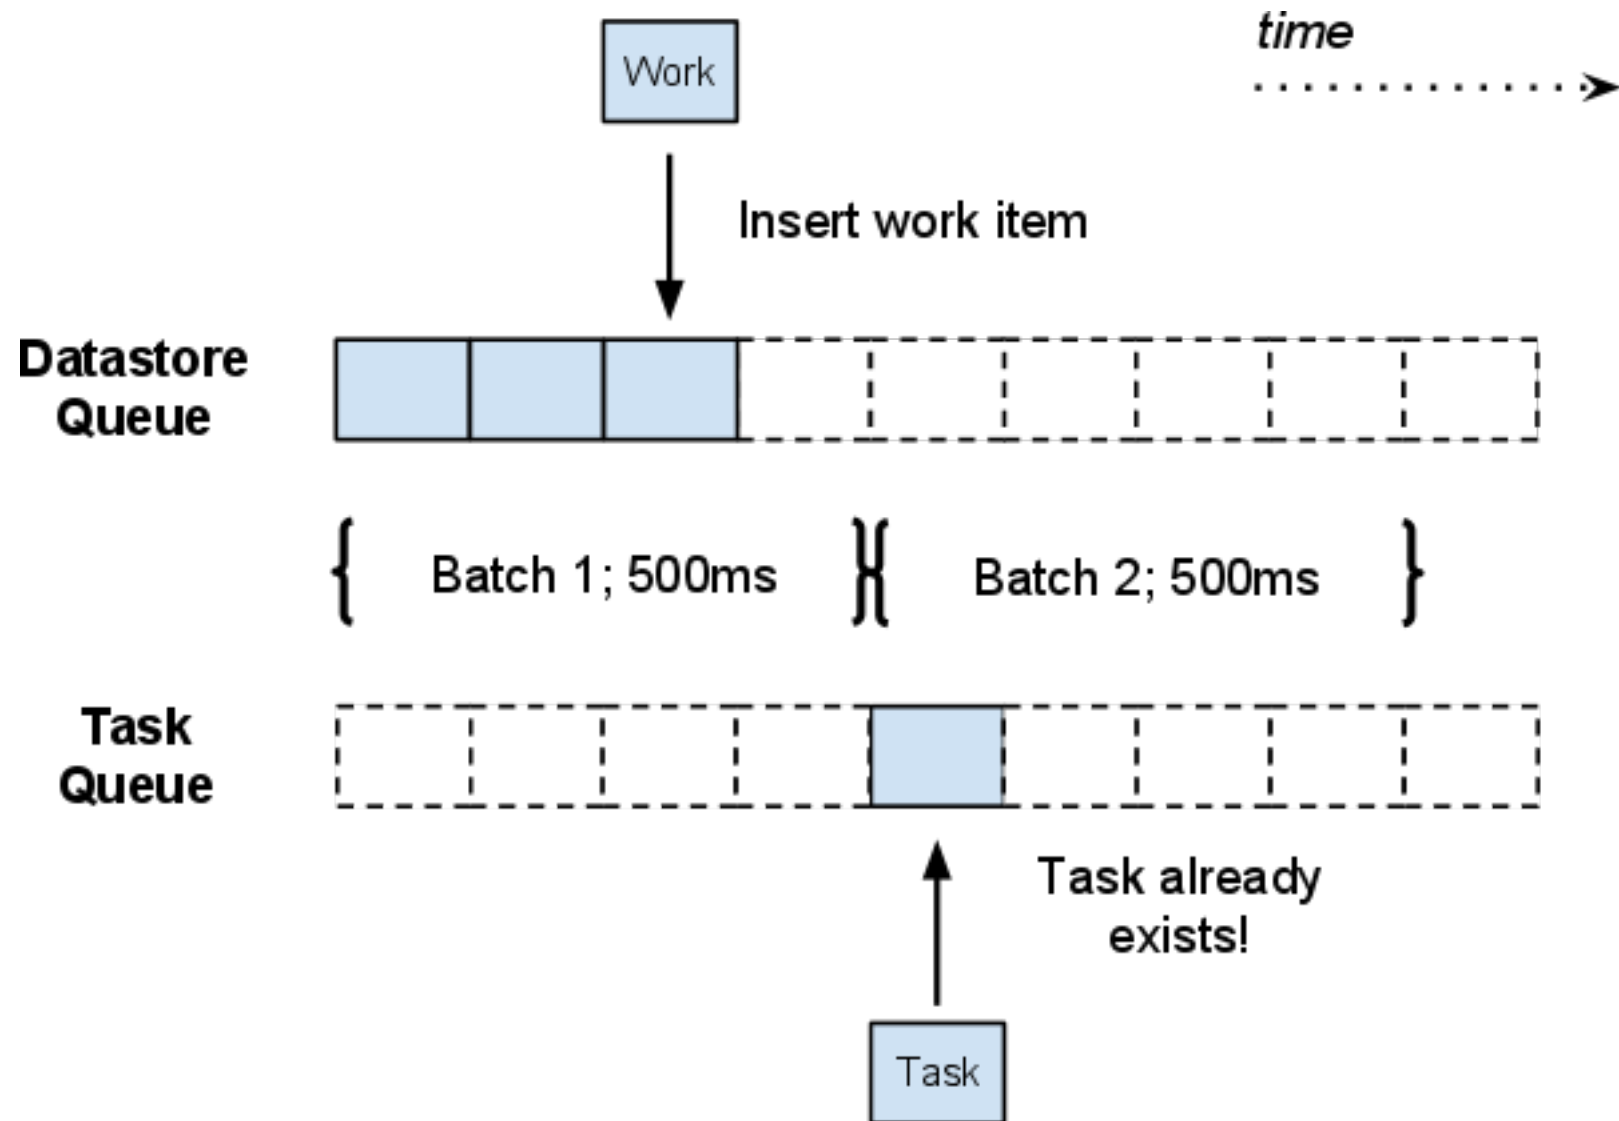
Fork-join queue with Datastore



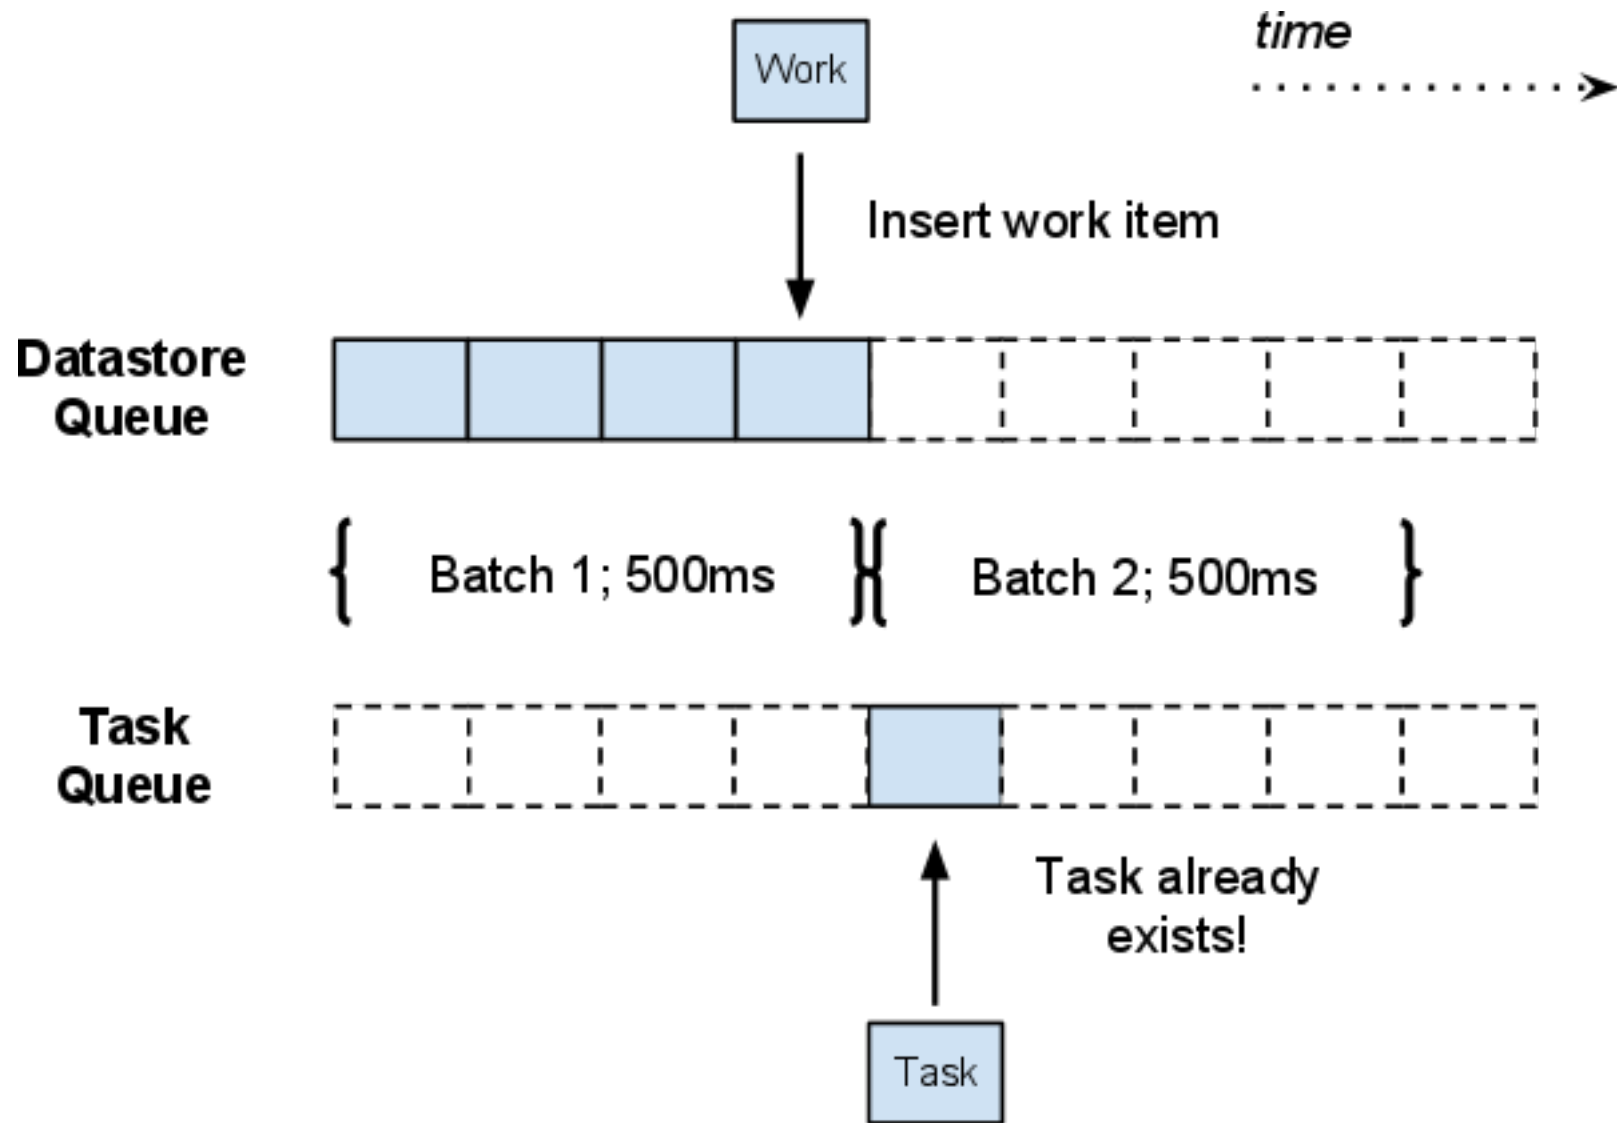
Fork-join queue with Datastore



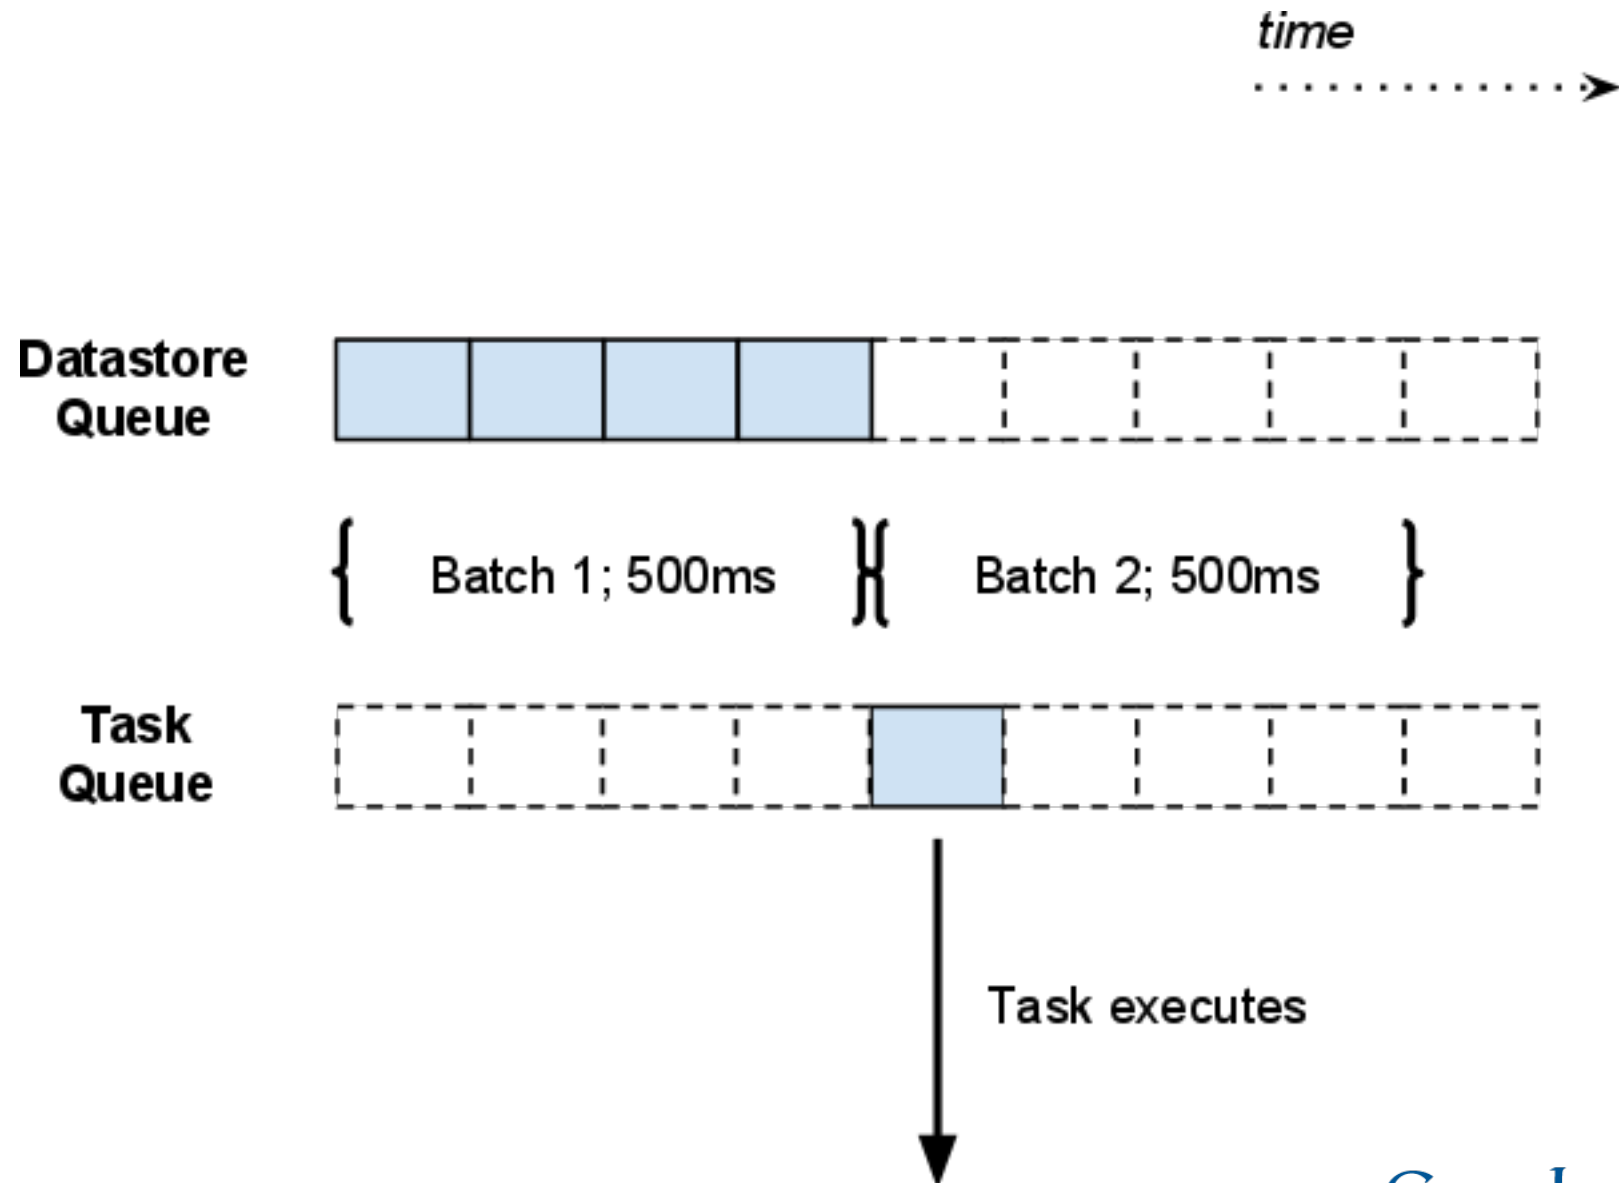
Fork-join queue with Datastore



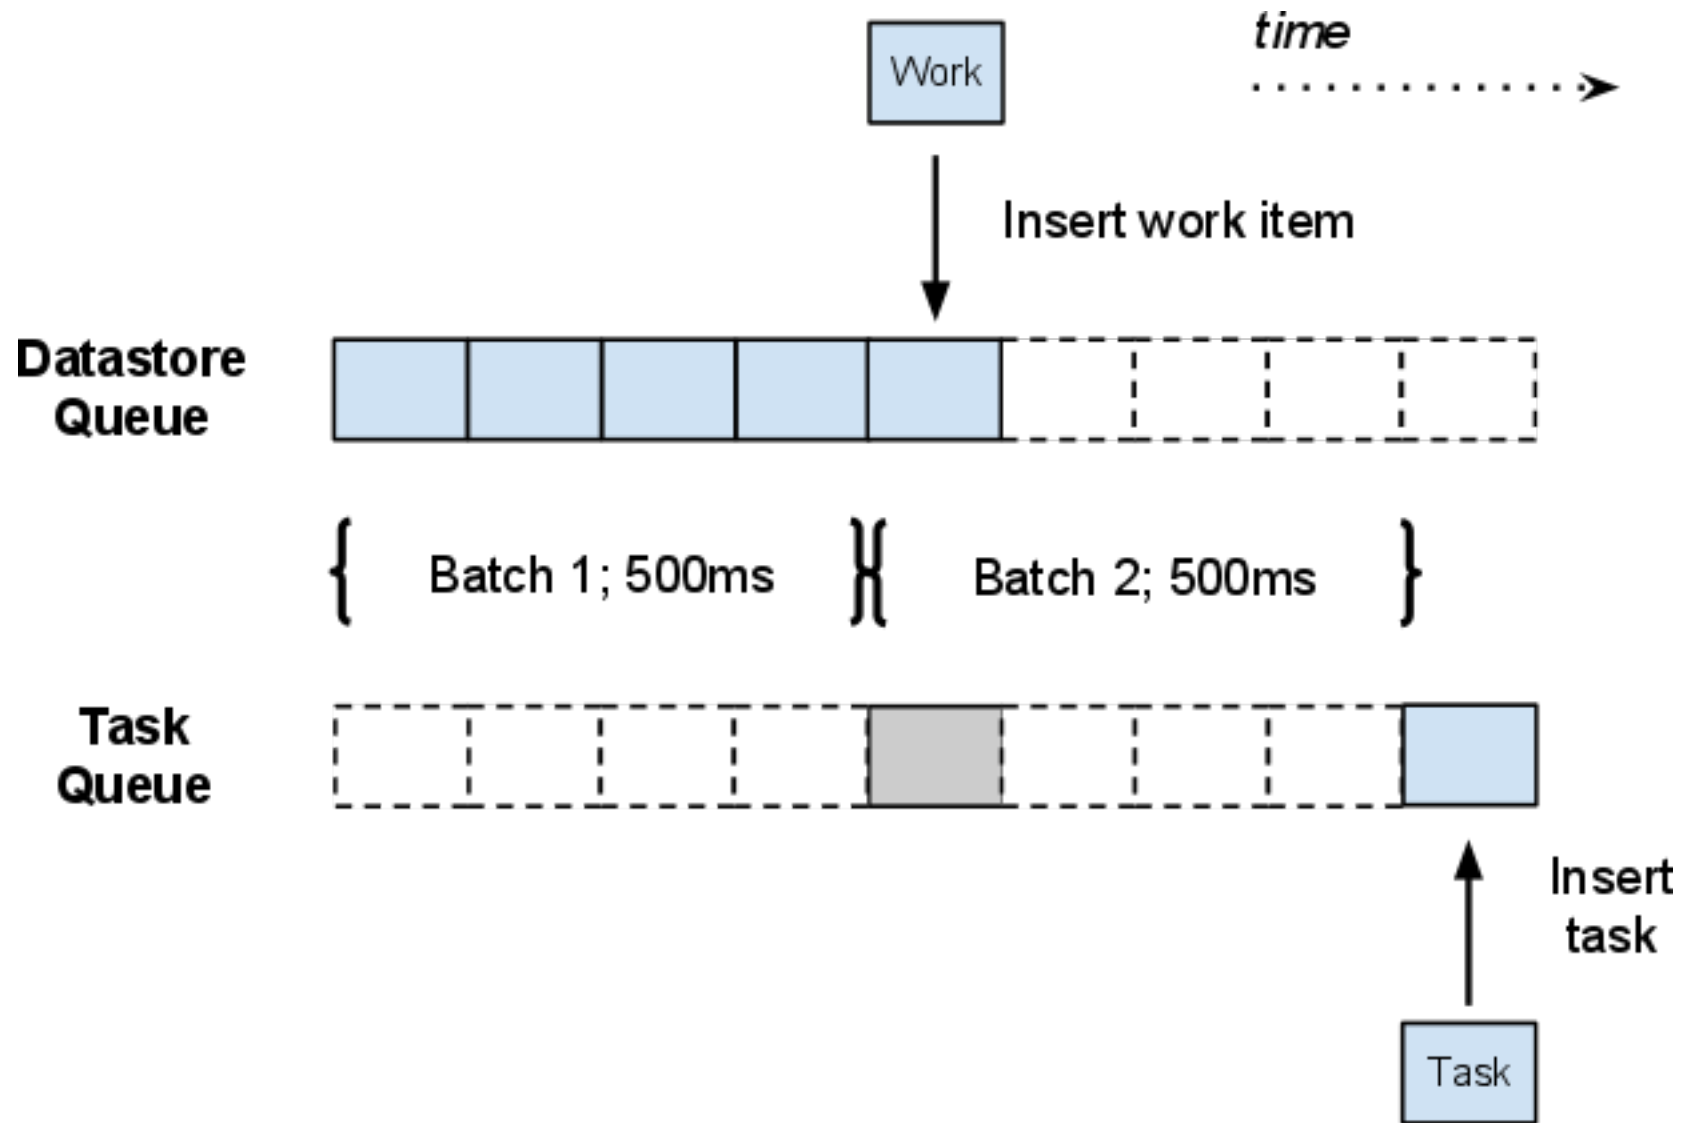
Fork-join queue with Datastore



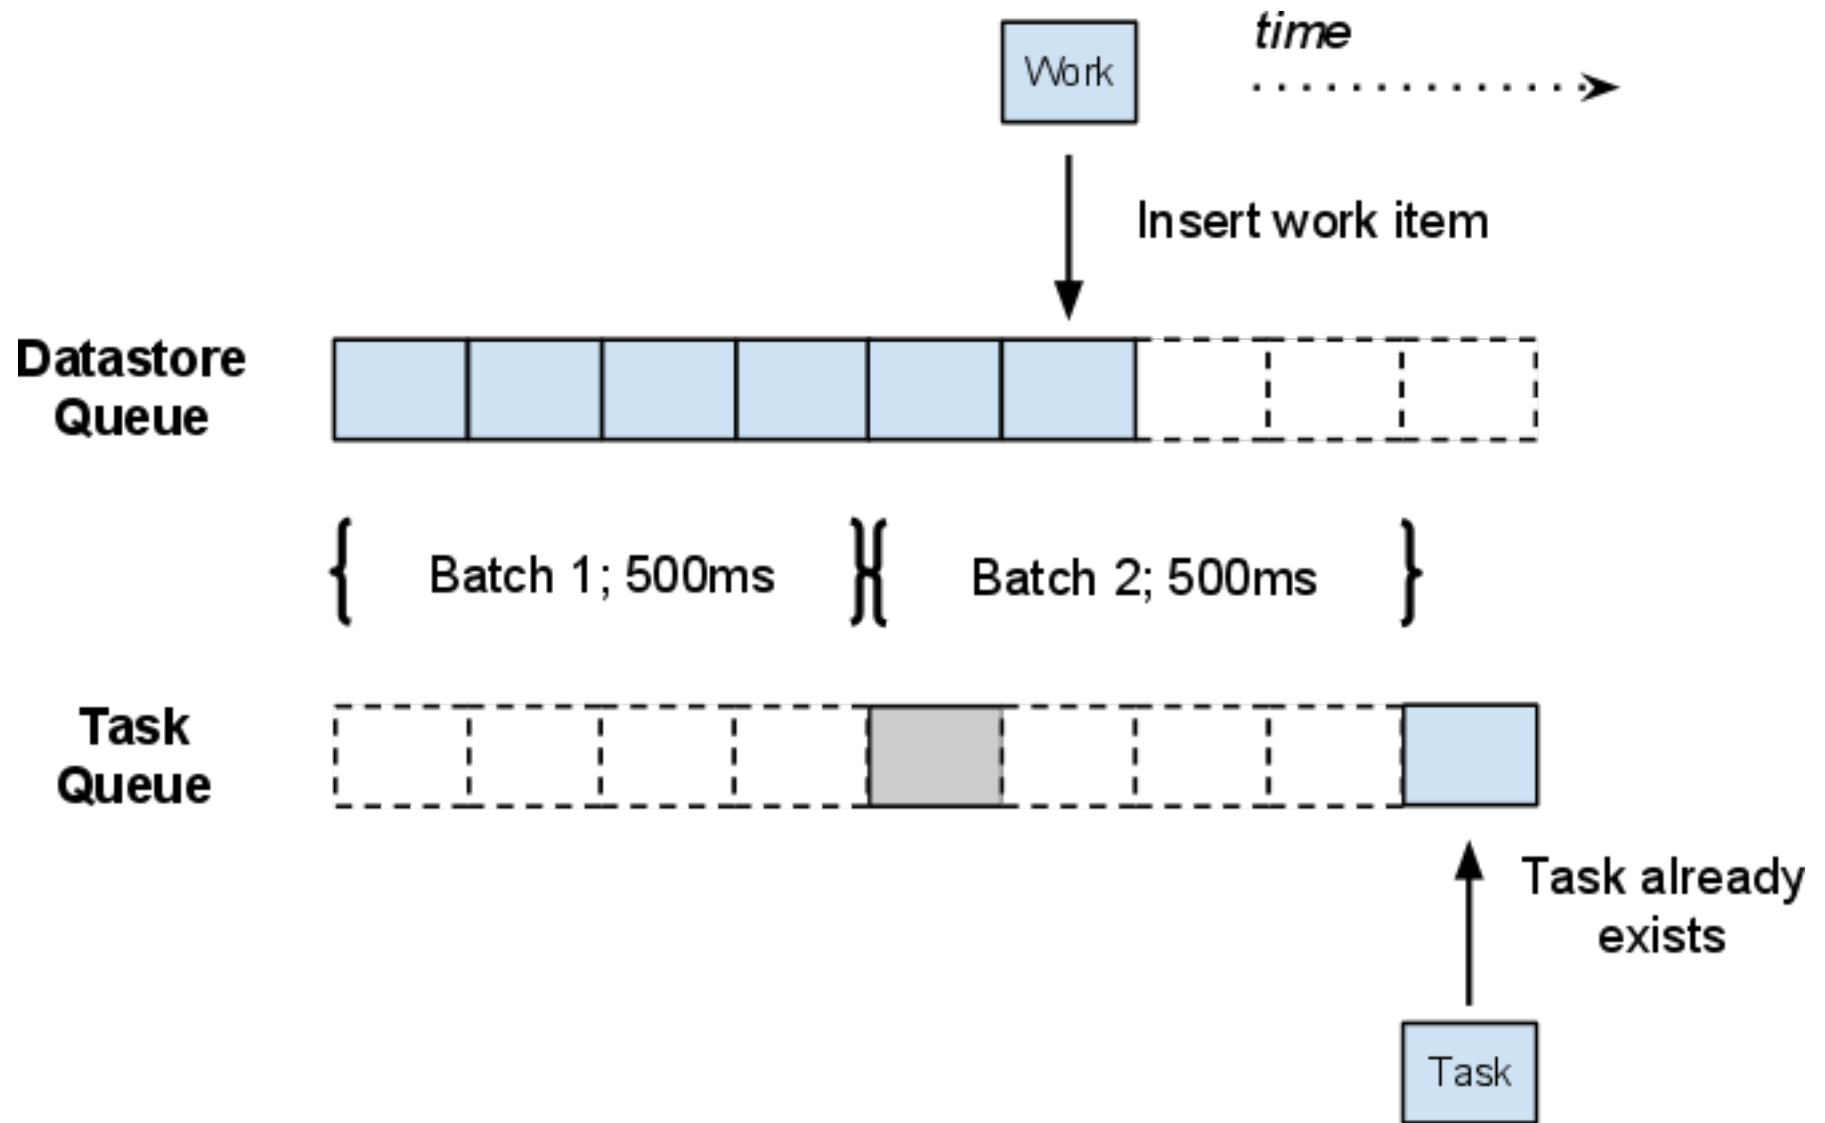
Fork-join queue with Datastore




Fork-join queue with Datastore



Fork-join queue with Datastore



Fork-join queue with Datastore: Race conditions



Fork-join queue example: Models

```
class MySum(db.Model):  
    name = db.StringProperty()  
    total = db.IntegerProperty()
```

```
class MyWork(db.Model):  
    work_index = db.StringProperty()  
    delta = db.IntegerProperty(indexed=False)
```

Fork-join queue example: Insert

```
def insert(sum_name, delta):
    index = memcache.get('index-' + sum_name)
    if index is None:
        memcache.add('index-' + sum_name, 1)
        index = memcache.get('index-' + sum_name)

    lock = '%s-lock-%d' % (sum_name, index)
    writers = memcache.incr(lock, initial_value=2**16)
    if writers < 2**16:
        memcache.decr(lock)
        return False # Insert fails, try again
    work = MyWork(delta=delta, work_index='%s-%d' %
                  (sum_name, knuth_hash(index)))
    work.put()
    # ... continues on next slide
```

Fork-join queue example: Insert continued

```
now = time.time()
try:
    taskqueue.add(
        name='%s-%d-%d' % (
            sum_name, int(now / 30), index),
        url='/work',
        eta=datetime.datetime.utcnow().timestamp(now) +
            datetime.timedelta(seconds=1))
except taskqueue.TaskAlreadyExistsError:
    pass # Fan-in magic
finally:
    memcache.decr(lock)

return True
```

Fork-join queue example: Join

```
def join(sum_name, index):  
    # force new writers to use the next index  
    memcache.incr('index-' + sum_name)  
  
    lock = '%s-lock-%d' % (sum_name, index)  
    memcache.decr(lock, 2**15) # You missed the boat  
  
    # busy wait for writers  
    for i in xrange(20): # timeout after 5s  
        counter = memcache.get(lock)  
        if counter is None or int(counter) <= 2**15:  
            break  
        time.sleep(0.250)  
  
    # ... continues on next slide
```

Fork-join queue example: Join continued

```
results = list(MyWork.all()  
    .filter('work_index =', '%s-%d' %  
        (sum_name, knuth_hash(index)))  
    .order('__key__'))  
delta = sum(r.delta for r in results)  
def txn():  
    my_sum = MySum.get_by_key_name(sum_name)  
    if my_sum is None:  
        my_sum = MySum(key_name=sum_name,  
            name=sum_name, total=0)  
    my_sum.total += delta  
    my_sum.put()  
db.run_in_transaction(txn)  
db.delete(results)
```

Fork-join queue example: Demo

Fork-join queue details

- Task names are the fan-in mechanism
- Task ETA for periodic batching
- memcache reader/writer locks for batch coordination
 - Spin locks with timeout
- Datastore queries to find work
- Use offline job to pick up drops (memcache failures)

Fork-join queue performance

- Depends on your batch size (work items per task)
 - Can achieve 80 to 1 easily.

Items per second	Average insert latency
22.1	223ms
39.7	245ms
55.6	258ms
75.8	249ms

Fork-join queue performance 2

- Work index **must** be a hash
 - Distribute load across Bigtable tablets
 - Alternative is tablet splits, unavailability
- Eliminate all other indexes on work items
 - Prevent overloading contiguous Bigtable rows
 - Can keep indexes if you're "boxcar"-ing transactions
- The magic of batch period = 0

Bonus: Fan-in with materialized views

Fan-in with materialized views: A sketch

1. Configure fan-in queue to batch once per second
2. User starts transaction on input data, update its value
 - Get fork-join work indexes for target aggregations
 - Assign work indexes to your input sequence markers
 - Enqueue update tasks (unnamed), Commit
3. Optimistically insert *named* fan-in tasks
 - Guarantees completion; ignored in common case
4. Later: Fan-in worker queries for inputs by work index
5. Worker transacts on aggregation data rows
 - Batch get of aggregation markers for inputs
 - Compare old markers to input sequence numbers
 - Compute commutative diff of up-to-date inputs
 - Update aggregation rows, Commit

Fan-in with materialized views: A sketch

- Please build this!

Future directions

Future directions

- Background servers
 - No wall-clock limits (30 sec deadline removed)
 - Chunk through fan-in queues in bulk
- Addressable servers
 - Send RPCs from user-facing requests to backends
 - Fan-in queues can be in memory
- Order of magnitude faster, skip disk writes

**View live notes and ask questions about
this session on Google Wave**

<http://tinyurl.com/app-engine-pipelines>

Me

<http://onebigfluke.com>

Google™

